

Are mental properties supervenient on brain properties?

Joshua T. Vogelstein¹, R. Jacob Vogelstein^{1,2}, Carey E. Priebe¹

¹Department of Applied Mathematics and Sciences,
Johns Hopkins University, Baltimore, MD, 21218,

²National Security Technology Department,
Johns Hopkins University Applied Physics Laboratory, Laurel, MD 20723

March 12, 2010

Abstract

The “mind-brain supervenience” theorem suggests that all mental properties (e.g. consciousness, intelligence, personality, etc.) are derived from the physical properties of the brain. The validity of this theorem has been argued in philosophical terms for over 2,500 years. **The alternative hypotheses, including various non-physical causes of mental properties, seems rather implausible to many. That said, testing these hypotheses** has not previously been approachable through experimental investigation. To address the question of whether the mind supervenes on the brain through empirical means, here we frame a supervenience hypothesis in rigorous mathematical terms and propose a modified version of supervenience (called ε -supervenience) that is amenable to scientific methods and statistical analysis. To elucidate this approach, we posit a thought experiment that illustrates how the probabilistic theory of pattern recognition can be used to make a one-sided determination of ε -supervenience. The physical property of the brain employed in this analysis is the graph describing brain connectivity (i.e., the *connectome*), and ε -supervenience allows us to determine whether a particular mental property can be inferred from one’s connectome to within any given misclassification rate $\varepsilon > 0$, regardless of the relationship between the two. In addition to the theoretical results, we show via simulation that given reasonable assumptions about class conditional probabilities and the amount of data available, the thought experiment can actually be conducted on a simple organism, *Caenorhabditis elegans*, with currently available technology.

Questioning the relationship between the mind (thoughts, beliefs, preferences, emotions, intelligence, etc.) and the brain (the physical structure inside our skulls) dates back at least as far as 400 BCE, when Plato wrote the dialogues,

in which he posited immateriality of the soul [1]. Approximately two millennia passed before these ideas reached their canonical form through Descartes’s discussion of mind-body dualism [2]. Then, in the 20th century, Donald Davidson stated and popularized the mind-brain supervenience theorem, which claims that an agent cannot alter in some mental property without altering in some physical property [3]. Contemporary fields of neural network theory and neuroscientific inquiry often assume mind-brain supervenience, or an even stronger assumption about mind-brain causality, but no previously proposed notion of supervenience seems amenable to empirical investigation. Here we define new versions of supervenience that formulate the theorem in rigorous mathematical terms and that can be experimentally tested.

The primary contributions of this work are as follows. First, a notion of supervenience emendable to empirical investigation is formally introduced. This renders the mind-brain dualism debate a hypothesis, rather than an assumption. Second, the limits of hypothesis testing to unravel supervenience hypotheses are provided. Third, we posit a very general model of brains and their associated mental properties that admits a universally consistent classifier; that is, we prove that assuming the brain can be accurately represented by this model, and there is sufficient data, this algorithm is guaranteed to obtain the best possible results. Fourth, simulations show this algorithm converges relatively rapidly.

Let \mathcal{B} be the observation space for some physical property, such as brain connectivity structure (i.e., connectome; see [4, 5, 6]). Let \mathcal{M} be the (finite) indicator space for some mental property, such as knowing calculus. Thus, for $b \in \mathcal{B}$ and $m \in \mathcal{M}$, the pair (b, m) represents a brain property/mind property pair.

Let $(B, M), (B_1, M_1), \dots, (B_n, M_n)$ be random observation pairs taking their values in $\mathcal{B} \times \mathcal{M}$, independently and identically distributed according to some joint probability distribution $F = F_{BM}$. Abusing notation to conceptually identify the properties with their spaces, the statistical supervenience relation $\mathcal{M} \stackrel{S}{\sim}_F \mathcal{B}$ says that $M_i \neq M_j \implies B_i \neq B_j$ (almost surely; where \implies does not suggest causation). That is, observing $B = b$ can allow us to assign m to M . While previously proposed notions of mind-brain-supervenience claim that all mental properties supervene on physical properties [7], here we consider empirically investigating only whether a particular mental property \mathcal{M} statistically supervenes on a particular physical property \mathcal{B} .

Let $g : \mathcal{B} \rightarrow \mathcal{M}$ be a classifier, which takes as input an observed brain connectivity structure b and produces a classification $\hat{m} = g(b)$ for the unobserved mental property m . The Bayes optimal classifier g^* minimizes $L_F(g)$ over all classifiers, where $L_F(g) = P_F[g(B) \neq M]$ denotes the probability of misclassification for classifier g under joint distribution $F = F_{BM}$. We can therefore rigorously define *statistical supervenience*:

Definition 1. \mathcal{M} is said to statistically supervene on \mathcal{B} for distribution $F = F_{BM}$, denoted $\mathcal{M} \stackrel{S}{\sim}_F \mathcal{B}$, if and only if $L_F(g^*) = 0$.

(Note that this definition does not imply a one-to-one mapping.) To allow

for the possibility of only *partial* supervenience,¹ we relax the above statistical supervenience to define ε -supervenience:

Definition 2. Given $\varepsilon > 0$, \mathcal{M} is said to ε -supervene on \mathcal{B} for distribution $F = F_{BM}$, denoted $\mathcal{M}^{\varepsilon}_F \mathcal{B}$, if and only if $L_F(g^*) < \varepsilon$.

Unfortunately, in general F is unknown, but can be estimated from the data. Therefore, generalizing the concept of a classifier g to allow consideration of training data, consider $g_n : \mathcal{B} \times (\mathcal{B} \times \mathcal{M})^n \rightarrow \mathcal{M}$ which takes as input an observed brain connectivity structure b and n training pairs $\vec{d}_n = (b_1, m_1), \dots, (b_n, m_n)$ and produces a classification $\hat{m} = g_n(b; \vec{d}_n)$. Let $L_F(g_n) = E[P_F[g_n(B; \vec{D}_n) \neq M | \vec{D}_n]]$.

Consider the problem of testing for ε -supervenience. Let the null hypothesis be given by $H_0 : L_F(g_n) \geq \varepsilon$ so that if we reject at level $\alpha > 0$ in favor of the alternative hypothesis $H_A : L_F(g_n) < \varepsilon$ then we can conclude, with $100(1 - \alpha)\%$ confidence, that $\mathcal{M}^{\varepsilon}_F \mathcal{B}$. Letting $\hat{L}_F^{n'}(g_n)$ denote the hold-out estimate of misclassification performance based on n' test observations, we note that $\hat{L}_F^{n'}(g_n)$ is distributed $Binomial(n', L_F(g_n))$. The test rejects for small $\hat{L}_F^{n'}(g_n)$. The level α critical value $c_\alpha(n', \varepsilon)$ is available under the least favorable distribution $Binomial(n', \varepsilon)$. Furthermore, $\mathcal{M}^{\varepsilon}_F \mathcal{B}$ implies $L_F(g^*) < \varepsilon$, and thus if g_n is a *consistent* classifier for $F = F_{BM}$ — that is, if $\lim_n L_F(g_n) = L_F(g^*)$ — then the power of this test (the probability of rejecting when in fact the alternative is true) goes to unity as $n, n' \rightarrow \infty$. Thus we have an inference procedure:

Theorem 1. Given $\alpha > 0$, we can test $\mathcal{M}^{\varepsilon}_F \mathcal{B}$ so that rejection implies $\mathcal{M}^{\varepsilon}_F \mathcal{B}$ holds with probability greater than or equal to $1 - \alpha$. Furthermore, given a consistent classifier the power of the test converges to unity.

Since the joint distribution $F = F_{BM}$ is unknown, the utility of Theorem 1 requires that g_n be a *universally consistent* classifier — that is, $\lim_n L_F(g_n) = L_F(g^*)$ for all distributions $F = F_{BM}$. Unfortunately, the rate of convergence of $L_F(g_n)$ to $L_F(g^*)$ depends on the (unknown) distribution $F = F_{BM}$ [8]. Furthermore, arbitrarily slow convergence theorems regarding the rate of convergence of $L_F(g_n)$ to $L_F(g^*)$ demonstrate that there is no universal n, n' which will guarantee that the test has power greater than any specified target $\beta > \alpha$ [9]. For this reason, the test outlined above can provide only a one-sided conclusion: if we reject we can be $100(1 - \alpha)\%$ confident that $\mathcal{M}^{\varepsilon}_F \mathcal{B}$ holds, but we can never be confident in its negation. Thus, without restrictions on F_{BM} , arbitrarily slow convergence theorems imply that our theorem of ε -supervenience does not *strictly* satisfy Popper's *falsifiability* requirement [10].² Given these caveats, consider the following thought experiment:

¹The following was stricken from the record, as I believed it to be a non-sequitor, but I partially replace it below: Unfortunately, it is in general impossible to determine whether $L_F(g^*) = 0$ without knowing F . Therefore,

²The strictly must be there because it is possible for the null to be rejected, given the data has a particular distribution, it is simply not necessarily rejectable. This is a subtle point that was lost on me until very recently.

Thought experiment 1. *Let the physical property under consideration be brain connectivity structure (“connectome”), so b is a graph (or, network) with vertices representing neurons (or neuroanatomical regions) and edges representing connections between neurons (or white matter tracts). Further let \mathcal{B} , the observation space, be the collection of all graphs on a finite number of vertices, and let $|\mathcal{B}|$ be countable. Now, imagine collecting very large amounts of very accurate independent and identically distributed brain-graph data and the associated mental property indicators. A k_n -nearest neighbor classifier using an isomorphism-matching Frobenius norm is universally consistent (see Appendix 1 for proof). Therefore, Theorem 1 applies and the existence of a universally consistent classifier guarantees that eventually (in n, n') we will be able to conclude $\mathcal{M}^{\varepsilon}_F \mathcal{B}$ for this mental/brain property pair, if indeed ε -supervenience holds. This logic holds for directed graphs or multigraphs or hypergraphs with discrete edge weights and vertex attributes.*

While the above thought experiment addresses the question of ε -supervenience, it does not address causality. Assuming we have confirmed $\mathcal{M}^{\varepsilon}_F \mathcal{B}$ for a particular mental/brain property pair with confidence level α , then morphing the brain (by altering edges) could be used to determine whether the relation is in fact causal.

Practical issues regarding actually conducting the above thought experiment include: (1) as stated, we must consider the space \mathcal{B} to be the quotient space of graphs mod graph isomorphism, unless the vertices are *labeled*; (2) a more informative and tractable distance on \mathcal{B} may be desired, as the k_n -nearest neighbor classifier under our Frobenius norm may have a rate of convergence so slow and a computational demand so high as to be impractical; and (3) collecting enough sufficiently accurate independent and identically distributed brain-graph data and the associated mental property indicators may be beyond current technological capabilities. Regardless, related experimental work includes collecting various types of brain graph data [11, 12, 13] and various approaches to inference on brain graphs [14, 15, 16], suggesting feasibility of such an experiment in the near future (see Appendix 2 for a simulated example of a feasible experiment). Nevertheless, our thought experiment suggests that we can hope to determine that a given mental property under consideration ε -supervenes on a brain’s connectivity structure. **This thought experiment, together with (i) the formal definition of ε -supervenience, (ii) the brain-graph model, and (iii) the universal consistency proof on graphs, is the first *proof* (to our knowledge) that “hope” of empirically investigating supervenience is possible.**

Appendix 1: k_n -nearest neighbor universal consistency for graphs

Assume first that all graphs are simple (meaning undirected with no loops and binary binary edges), on the same set of vertices, and that the graphs are labeled so that we know which vertex in one graph corresponds to which vertex in

another. Then the Frobenius distance function $d(b_1, b_2)$ can be written in terms of the associated adjacency matrices A_1 and A_2 : $d(b_1, b_2) = \|A_1 - A_2\|_F$. If the graphs are identical, then $d(b_1, b_2) = 0$, and if the graphs are different, then $d(b_1, b_2) \geq 1$. Since the space \mathcal{B} is finite, n large enough guarantees that with probability approaching unity at least k_n training samples coincide with each atom, so long as $k_n/n \rightarrow 0$. Then $k_n \rightarrow \infty$ guarantees that the nearest neighbor vote-winner for each atom will eventually coincide with Bayes' choice, yielding universal consistency.

In the foregoing argument, there exists a smallest non-zero atomic probability p_{min} , and “ n large enough” is driven by this probability. Generalizing to countable \mathcal{B} with discrete weights, we see that given $\delta > 0$, there is a finite set S with $P[S] > 1 - \delta$ and smallest atomic probability p_{min} , so that $L_F(g_n) \rightarrow c \leq L_F(g^*) + \delta$, yielding universal consistency.

If the graphs may have different numbers of vertices, and are unlabeled, we consider the isomorphism-matching Frobenius norm. Assume without loss of generality that b_1 has at least as many vertices as b_2 , and write A_2^P for the adjacency matrix associated with b_2 “padded” to include extra isolated vertices so that A_2^P is the same size as A_1 . Then $d(b_1, b_2) = \min_Q \|QA_1Q^T - A_2^P\|_F$ where the minimum is taken over all permutation matrices [17]. Under the equivalence relation induced by this isomorphism-matching, the foregoing universal consistency argument holds.

Several points of note: isolated vertices are ignored in our equivalence relation; the class-conditional signal is entirely encompassed by the connectivity structure; the graph isomorphism problem is computationally hard [18, 19]; and the argument employed here does not capture the concept of “nearness implies likelihood of similar class”—we simply rely on atomic behavior.

Appendix 2: Simulation

As an example of a feasible experiment, one may consider a species whose nervous system consists of the same (small) number of labeled neurons for each organism. *Caenorhabditis elegans* is believed to be such a species [20]. The hermaphroditic *C. elegans*' somatic nervous system consists of 279 interconnected neurons. While the graph with these neurons as vertices and edges defined by chemical synapses between neurons is not identical across individuals, it is reasonably consistent [20]. Furthermore, these animals exhibit a rich behavioral repertoire that depends on circuit properties [21]. Thus, one may design an experiment by describing the joint distribution F_{BM} via class-conditional distributions $F_{B|M=m_j}$ for the *C. elegans* brain-graph for two mental properties of interest, m_0 and m_1 , along with the prior probability of class membership $P[M = m_1]$. Here the mental property corresponds to the *C. elegans* exhibiting or not exhibiting a particular behavior (e.g., response to an odor).

Simulations suggest that one may build a classifier, practically and with a manageable training sample size n , that demonstrates ε -supervenience with reasonable choices for ε and α and a plausible joint distribution F_{BM} (Figure 1).

To generate the data, we let the class-conditional random variable $E_{ij}|M = m_0$ be distributed $\text{Poisson}(A_{ij} + \eta)$, where A_{ij} is the number of chemical synapses between neuron i and neuron j according to [22], with noise parameter $0 < \eta \ll 1$. The class-conditional random variable $E_{ij}|M = m_1$ is distributed $\text{Poisson}(A_{ij} + z_{ij})$ for neurons $i, j \in \mathcal{D}$, where \mathcal{D} is the set of edges deemed responsible for odor-evoked behavior according to [23], with signal parameter z_{ij} uniformly sampled from $[-5, 5]$. We consider k_n -nearest neighbor classification of labeled multigraphs (directed, with loops) on 279 vertices, under Frobenius norm. The k_n -nearest neighbor classifier used here satisfies $k_n \rightarrow \infty$ as $n \rightarrow \infty$ and $k_n/n \rightarrow 0$ as $n \rightarrow \infty$, ensuring universal consistency. (Better classifiers can be constructed for the joint distribution F_{BM} used here; however, we demand universal consistency.)

Importantly, conducting this experiment *in actu* is not beyond current technological limitations. 3D superresolution imaging [24] combined with neurite tracing algorithms [25, 15, 16] allow the collection of a brain-graph within a day. Genetic manipulations, laser ablations, and training paradigms can each be used to obtain a non-wild type population for use as $M = m_1$ [21], and the class of each organism (m_0 vs. m_1) can also be determined automatically [26].

Acknowledgments The authors would like to acknowledge helpful discussions with J Lande and B Vogelstein. This work was supported in part by the NSA Research Program in Applied Neuroscience.

References

- [1] Plato, *Plato: complete works*. Hackett Pub Co, 1997.
- [2] R. Descartes, *Meditationes de prima philosophia*. 1641.
- [3] D. Davidson, *Experience and Theory*, ch. Mental Events. Duckworth, 1970.
- [4] O. Sporns, G. Tononi, and R. Kotter, “The human connectome: A structural description of the human brain,” *PLoS Computational Biology*, vol. 1, no. 4, p. e42, 2005.
- [5] J. W. Lichtman, J. Livet, and J. R. Sanes, “A technicolour approach to the connectome,” *Nat Rev Neurosci*, vol. 9, pp. 417–422, Jun 2008.
- [6] H. Seung, “Reading the Book of Memory: Sparse Sampling versus Dense Mapping of Connectomes,” *Neuron*, vol. 62, no. 1, pp. 17–29, 2009.
- [7] J. Kim, *Philosophy of Mind*. Westview Press, second ed., 2005.
- [8] L. Devroye, L. Györfi, and G. Lugosi, *A Probabilistic Theory of Pattern Recognition*. Springer, 1996.
- [9] L. Devroye, “On arbitrarily slow rates of global convergence in density estimation,” *Probability Theory and Related Fields*, vol. 62, no. 4, pp. 475–483, 1983.

- [10] K. Popper, “The logic of scientific discovery,” 1959.
- [11] J. White, E. Southgate, J. N. Thomson, and S. Brenner, “The structure of the nervous system of the nematode *caenorhabditis elegans*,” *Philosophical Transactions of Royal Society London. Series B, Biological Sciences*, vol. 314, no. 1165, pp. 1–340, 1986.
- [12] W. W. Denk and H. Horstmann, “Serial block-face scanning electron microscopy to reconstruct three-dimensional tissue nanostructure,” *PLOS Biol.*, vol. 2, p. e329, 2004.
- [13] K. Briggman and W. Denk, “Towards neural circuit reconstruction with volume electron microscopy techniques,” *Current opinion in neurobiology*, vol. 16, no. 5, pp. 562–570, 2006.
- [14] J. H. Macke, N. Maack, R. Gupta, W. Denk, B. Schlkopf, and A. Borst, “Contour-propagation algorithms for semi-automated reconstruction of neural processes,” *J Neurosci Methods*, vol. 167, pp. 349–357, Jan 2008.
- [15] Y. Mishchenko, “Automation of 3d reconstruction of neural tissue from large volume of conventional serial section transmission electron micrographs,” *J Neurosci Methods*, vol. 176, pp. 276–289, Jan 2009.
- [16] J. Lu, J. C. Fiala, and J. W. Lichtman, “Semi-automated reconstruction of neural processes from large numbers of fluorescence images,” *PLoS ONE*, vol. 4, p. e5655, 05 2009.
- [17] R. Horn and C. Johnson, *Matrix analysis*. Cambridge Univ Pr, 1990.
- [18] J. M. Conroy, S. G. Kratzer, and L. J. Podrazik, “A continuous method for the quadratic assignment problem,” in *Society for Industrial and Applied Mathematics*, 1997.
- [19] M. Zaslavskiy, F. Bach, and J. Vert, “A path following algorithm for the graph matching problem,” *IEEE Transactions on Pattern Analysis and Machine Intelligence*, vol. 31, pp. 2227–2242, 2008.
- [20] R. M. Durbin, *Studies on the Development and Organisation of the Nervous System of Caenorhabditis elegans*. PhD thesis, University of Cambridge, 1987.
- [21] M. de Bono and A. V. Maricq, “Neuronal substrates of complex behaviors in *c. elegans*,” *Annu Rev Neurosci*, vol. 28, pp. 451–501, 2005.
- [22] L. Varshney, B. Chen, E. Paniagua, D. Hall, and D. Chklovskii, “Structural Properties of the *Caenorhabditis elegans* Neuronal Network,” *ArXiv*, 2009.
- [23] S. H. Chalasani, N. Chronis, M. Tsunozaki, J. M. Gray, D. Ramot, M. B. Goodman, and C. I. Bargmann, “Dissecting a circuit for olfactory behaviour in *caenorhabditis elegans*,” *Nature*, vol. 450, pp. 63–70, Nov 2007.

- [24] A. Vaziri, J. Tang, H. Shroff, and C. V. Shank, “Multilayer three-dimensional super resolution imaging of thick biological samples.,” *Proc Natl Acad Sci U S A*, vol. 105, pp. 20221–20226, Dec 2008.
- [25] M. Helmstaedter, K. L. Briggman, and W. Denk, “3d structural imaging of the brain with photons and electrons.,” *Curr Opin Neurobiol*, vol. 18, pp. 633–641, Dec 2008.
- [26] S. D. Buckingham and D. B. Sattelle, “Strategies for automated analysis of *c. elegans* locomotion.,” *Invert Neurosci*, vol. 8, pp. 121–131, Sep 2008.




Figure 1: *C. elegans* graph classification simulation results. $\hat{L}_F^{1000}(g_n)$ is plotted as a function of class-conditional training sample size n_j , suggesting that for $\varepsilon = 0.1$ we can determine that $\mathcal{M}_{\tilde{F}\mathcal{B}}^\varepsilon$ holds with 99% confidence with just a few hundred training samples generated from F_{BM} . Each dot depicts an estimate for $L_F(g_n)$; standard errors are $(L_F(g_n)(1 - L_F(g_n))/1000)^{1/2}$; e.g., $n_j = 180$; $k_n = 53$; $\hat{L}_F^{1000}(g_n) = 0.057$; standard error less than 0.01. We reject $H_0 : L_F(g^*) \geq 0.10$ at $\alpha = 0.01$. $L_F(g^*) \approx 0$ for this simulation.